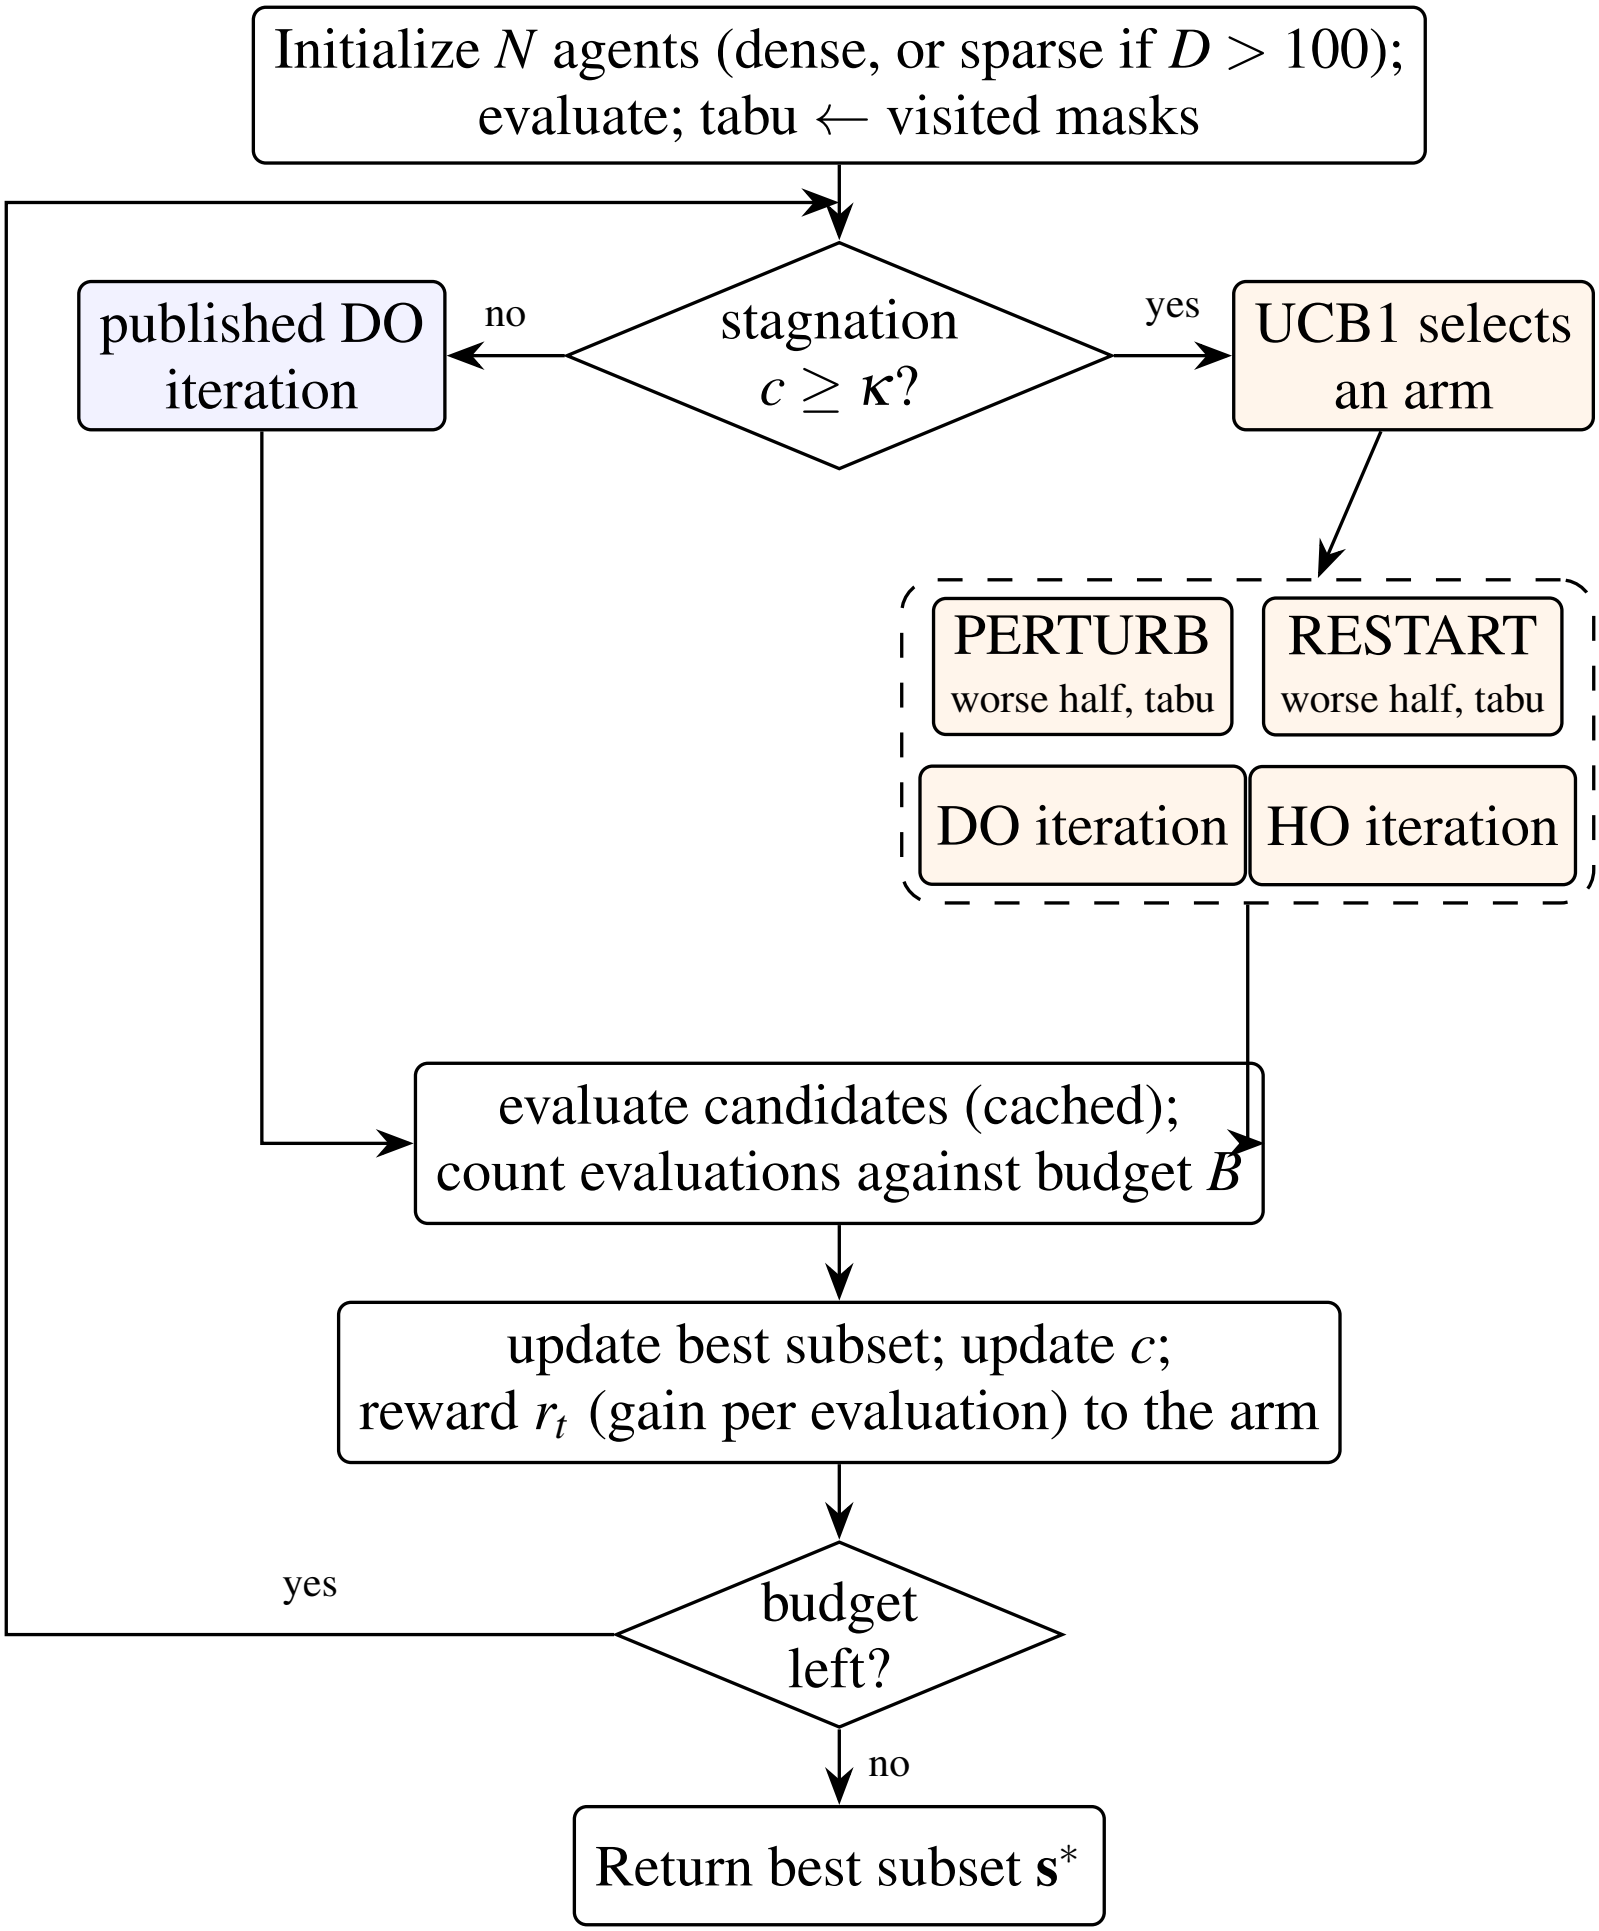
Initialize N agents (dense, or sparse if $D > 100$); evaluate; tabu ← visited masks
stagnation $c \geq \kappa$?
no
yes
published DO iteration
UCB1 selects an arm
PERTURB
worse half, tabu
RESTART
worse half, tabu
DO iteration
HO iteration
evaluate candidates (cached); count evaluations against budget B
update best subset; update c; reward r_t (gain per evaluation) to the arm
budget left?
yes
no
Return best subset $\mathbf{s}^*$